

Yassine Kraiem

New York, NY | yassinekraiem08@gmail.com | +1 (808) 570-8151 | linkedin.com/in/yassinekraiem | yassinekraiem.com

EDUCATION

Columbia University

M.S. in Artificial Intelligence.

New York, NY

Dec 2027

Alma College

B.S. in Computer Science. Summa Cum Laude. GPA: 3.88.

Alma, MI

Dec 2025

EXPERIENCE

Founder and AI Software Engineer

Jul 2025 – Jan 2026

RecipeOne

Alma, MI

- Built a multi-tenant B2B SaaS platform for institutional foodservice operations, including recipe planning, procurement, and inventory management across multisite facilities, with a PostgreSQL backend modeling 15+ entities and 30+ relational tables.
- Implemented automated unit conversion and yield computation pipelines, processing 7,000+ ingredient records to generate purchasing quantities and cost estimates across inventory systems.

Co-Founder and AI Software Engineer

Oct 2023 – Apr 2025

ListFlowAI

Remote, Washington D.C.

- Co-founded an AI-powered B2B SaaS automation platform used by 200+ users, processing 5,000+ inbound leads through LLM-driven enrichment, classification, and routing workflows.
- Architected backend services and data pipelines using Python, FastAPI, and PostgreSQL, extracting structured data from emails, forms, and CRM exports and reducing manual lead processing by 80%.

Software Engineering Intern

Jun 2024 – Aug 2024

Eat Pro Japan

Tokyo, Japan

- Shipped features for a booking and reservation system (Node.js, Next.js, MySQL on AWS) on a 20-engineer team, hardening reservation flow against availability conflicts as the platform scaled from 5K to 20K users.
- Built an internal admin tool for the operations team to manage restaurant availability and resolve booking conflicts in real time, replacing a manual spreadsheet-based workflow and eliminating a recurring class of user-facing reservation failures.

PROJECTS

AI Decision Support System (Retrieval-Augmented Generation)

Oct 2025 – Feb 2026

<https://github.com/Yaskm08/rag-decision-support-system>

- Developed a production RAG system using Python, FastAPI, and pgvector, enabling hybrid retrieval across 1,600+ document chunks and reducing hallucinated responses from 80% to 0% through citation-grounded generation.
- Designed a parallel reranking and semantic caching pipeline, cutting p95 latency by 76%; conducted a confidence calibration study (ECE = 0.256) and documented 9 failure modes via CI/CD evals.

Natural Language Analytics System (Text-to-SQL)

Nov 2025 – Mar 2026

<https://github.com/Yassinekraiem08/nl-to-sql-analytics-system>

- Developed a full-stack NL-to-SQL analytics system (FastAPI, React, SSE streaming), adapting ideas from DIN-SQL and CHESSE for execution-guided refinement and schema-aware generation over complex relational schemas.
- Engineered a self-correcting SQL pipeline with up to 3 retry attempts, a semantic query cache (cosine similarity ≥ 0.92 , 0 LLM calls on hits), Spider-style execution-accuracy benchmarking with hardness classification, and 212 automated tests across 15 modules.

AI Workflow Orchestrator (Agents + Tool Calling)

Nov 2025 – Mar 2026

<https://github.com/Yassinekraiem08/ai-workflow-orchestrator>

- Built a multi-agent AI orchestration system for incident triage across logs, tickets, and emails using a classify, plan, execute, replan loop with human-in-the-loop escalation; deployed on AWS ECS Fargate with observability via Jaeger, Prometheus, Grafana, and CI/CD through GitHub Actions OIDC.
- Designed fault-tolerant execution with Celery workers, retries with exponential backoff, and dead-letter queues; implemented multi-model routing and LLM-as-judge evaluation across 5 metrics, achieving 95% task success vs. 68% baseline and reducing per-task cost by 21%.

SKILLS

AI/ML: LLM Systems, RAG, Agent Orchestration, PyTorch, Core ML, Model Evaluation

Languages: Python, C++, Java, Go, Swift, TypeScript, SQL

Backend: FastAPI, gRPC, Protocol Buffers, REST APIs, Microservices

Systems: Distributed Systems, Concurrency, Data Structures and Algorithms, System Design

Infrastructure: Linux, Docker, Kubernetes, AWS (ECS, Lambda, S3), CI/CD (GitHub Actions)

Databases: PostgreSQL, Redis, Vector Search