

YASSINE KRAIEM

New York, NY | +1 (808) 570-8151 | yassinekraiem08@gmail.com | linkedin.com/in/yassinekraiem | github.com/Yassinekraiem08

EDUCATION

Columbia University, New York
M.S. in Artificial Intelligence

Aug 2026 — May 2028

Alma College, Alma
B.S. in Computer Science

Aug 2022 — Dec 2025
GPA: 3.88, Summa Cum Laude

EXPERIENCE

RecipeOne

Jul 2025 — Jan 2026

Founder & AI Software Engineer

Alma, MI

- Built a multi-tenant B2B SaaS platform for institutional foodservice ops (recipe planning, procurement, inventory across multi-site facilities), with PostgreSQL backend modeling 15+ entities and 30+ relational tables.
- Implemented automated unit conversion and yield computation pipelines, processing 7,000+ ingredient records to generate purchasing quantities and cost estimates across inventory systems.

ListFlowAI — <https://listflowai.com>

Oct 2023 — Apr 2025

Co-Founder & AI Software Engineer

Remote - Washington, D.C.

- Co-founded an AI-powered B2B SaaS automation platform used by 200+ users, processing 5,000+ inbound leads through LLM-driven enrichment, classification, and routing workflows.
- Architected backend services and data pipelines using Python, FastAPI, and PostgreSQL, extracting structured data from emails, forms, and CRM exports and reducing manual lead processing by 80%.

Kinney Research Fellow — Autonomous Drone Delivery Simulation Platform

Feb 2025 — Jul 2025

Lead Software Engineer (Full-stack)

Alma, MI

- Designed and built a full-stack simulation platform modeling autonomous drone delivery under dynamic weather and demand, executing 1,000+ simulated delivery scenarios for system evaluation and performance benchmarking.
- Engineered route optimization and scheduling algorithms using graph-based and heuristic methods, improving delivery time by 18% and reducing energy consumption by 12% across fleet operations.

PROJECTS

AI Decision Support System (Retrieval-Augmented Generation)

Oct 2025 — Feb 2026

- Built a production RAG system using Python, FastAPI, and pgvector, enabling hybrid retrieval across 1,600+ document chunks and reducing hallucinated responses from 80% to 0% through citation-grounded generation.
- Designed parallel reranking and semantic caching pipeline, cutting p95 latency by 76%; conducted confidence calibration study (ECE = 0.256) and documented 9 failure modes via CI/CD evals.

Natural Language Analytics System (Text-to-SQL) | querymind-demo.vercel.app

Nov 2025 — Mar 2026

- Built a full-stack NL-to-SQL analytics system (FastAPI, React, SSE streaming), adapting ideas from DIN-SQL and CHESS for execution-guided refinement and schema-aware generation over complex relational schemas.
- Engineered a self-correcting SQL pipeline with up to 3 retry attempts, semantic query cache (cosine similarity ≥ 0.92 , 0 LLM calls on hits), Spider-style execution-accuracy benchmark with hardness classification, and 212 automated tests across 15 modules.

AI Workflow Orchestrator (Agents + Tool Calling) | <https://ai-workflow-orchestrator.vercel.app>

Nov 2025 — Mar 2026

- Built a multi-agent AI orchestration system for incident triage across logs, tickets, and emails using a classify → plan → execute → replan loop with human-in-the-loop escalation; deployed on AWS ECS Fargate with observability via Jaeger, Prometheus, Grafana, and CI/CD through GitHub Actions OIDC.
- Designed fault-tolerant execution with Celery workers, retries with exponential backoff, and dead-letter queues; implemented multi-model routing and LLM-as-judge evaluation (5 metrics), achieving 95% task success vs 68% baseline and reducing per-task cost by 21%.

SKILLS

Languages: Python, C++, Java, Go, Swift, TypeScript, SQL

Systems: Distributed Systems, Concurrency, Data Structures & Algorithms, System Design

Backend: gRPC, Protocol Buffers, REST APIs, Microservices, FastAPI

Infrastructure: Linux, Docker, Kubernetes, AWS (ECS, Lambda, S3), CI/CD (GitHub Actions)

Databases: PostgreSQL, Redis, Vector Search

AI/ML: LLM Systems, RAG, Agent Orchestration, PyTorch, Core ML, Model Evaluation